

📌 IoT Data Analytics Architecture – Portfolio Case Study

📌 Project Overview

Designed and documented a scalable, secure, and cost-effective IoT data processing architecture on AWS for a **transport and payments solutions use case**.

The solution ingests real-time data from **point-of-sale (POS) devices and ferry terminals**, processes and transforms the data, and provides actionable insights for executive decision-making, operations optimisation, and machine learning analytics.

📌 Problem Statement

The organisation needed to:

- Capture real-time data streams from IoT POS and ferry devices.
 - Store raw data for compliance and ML training.
 - Transform data into analytics-ready formats.
 - Provide dashboards to operations, finance, and executive teams for quick data-driven decisions.
 - Maintain a scalable, highly available, and secure architecture while optimising cost.
-

📌 Tools & Technologies

- **AWS API Gateway** – Secure ingestion of POST requests from IoT devices.
 - **AWS Lambda** – Serverless data processing and routing logic.
 - **AWS Glue & Glue Crawler** – Data transformation and schema cataloguing.
 - **Amazon S3** – Storage for raw and transformed data.
 - **Amazon DynamoDB** – Metadata and index storage for quick retrieval and tracking.
 - **Amazon QuickSight** – Dashboards and business intelligence.
 - **Amazon SageMaker** – Machine learning model training using historical data.
 - **IAM Roles & Policies** – Fine-grained access controls and security.
-

⚙️ Solution Architecture

- 1. Data Ingestion**
IoT devices send POST requests → API Gateway → triggers Lambda.
 - 2. Processing & Storage**
Lambda stores raw data in S3 raw bucket, writes metadata to DynamoDB, and triggers Glue jobs.
 - 3. Data Transformation**
Glue Jobs process data → store cleaned/transformed output in S3 transformed bucket → Glue Crawler updates schema in Data Catalog.
 - 4. Analytics & Machine Learning**
QuickSight connects to DynamoDB and transformed S3 data for dashboards.
SageMaker accesses raw data in S3 for ML training, referencing DynamoDB for dataset tracking.
 - 5. Security & Scalability**
IAM roles with least privilege for Lambda, Glue, QuickSight, and SageMaker.
Serverless and multi-AZ services ensure high availability and cost efficiency.
-

📌 Challenges Faced

- Designing for high throughput ingestion with low latency.
 - Managing schema evolution and continuous cataloguing.
 - Ensuring secure access controls across multiple services.
 - Architecting for cost optimisation while maintaining performance and scalability.
-

📌 Solutions Implemented

- Used **AWS Glue Crawler** to dynamically catalog evolving datasets.
- Designed Lambda functions with efficient error handling and retry logic for ingestion stability.
- Integrated **DynamoDB** as a lightweight metadata index for rapid dataset tracking and referencing.
- Configured **IAM roles and policies** based on least privilege for security compliance.
- Proposed **serverless-first architecture** (API Gateway + Lambda + Glue) to eliminate idle resource costs.

📌 Impact & Results

- Enabled real-time data availability for executive and operations dashboards.
 - Reduced operational complexity through automation with Glue and Lambda.
 - Improved data governance via Glue Data Catalog and DynamoDB metadata indexing.
 - Laid the foundation for advanced machine learning use cases with SageMaker.
 - Designed with multi-zonal high availability, ensuring zero downtime for critical ingestion pipelines.
-

👤👤👤 My Role

- Architected and documented the end-to-end solution.
 - Designed IAM policies, data flows, and security considerations.
 - Proposed implementation patterns for Lambda and Glue integration.
 - Provided scalability and cost optimisation recommendations.
 - Developed detailed technical documentation for stakeholder presentations.
-

📌 Key Takeaways

- 📌 Serverless architectures simplify scaling and reduce costs.
 - 📌 DynamoDB is effective for lightweight metadata and indexing.
 - 📌 Glue Crawlers automate schema management, enabling seamless analytics.
 - 📌 Clear IAM role boundaries are critical for security and audit compliance.
-

📌 Cost Considerations

This architecture is **documented without live AWS resource deployment** in this repository to avoid incurring costs. It serves as a detailed reference and portfolio showcase for implementation in production environments.
